

Same cell, two runs: 62.0
and 47.2. Five calls each.

62.0

RUN ONE,
MEAN OF 5

47.2

RUN TWO, MEAN OF 5

Noise 3.0 and 4.8 points. Five reps run.

Same posting, model and resume, two runs 90 minutes apart. Run-to-run SD per cell: 3.0 and 4.8.

One resume, 21 postings from 8 employers, cheap models, a research prompt. Under 3 points: not excluded.

One model repeated itself. The other never did.

haiku-4-5		17
	17 of 42 cells identical	
gpt-5-mini		0
	0 of 42	

Identical scores on 17 of 42 cells, and 0 of 42.

Cells where every rep returned the same score. Standard deviation per cell: 3.0 and 4.8 points.

claude-haiku-4-5 returned 10 distinct scores in the whole run. gpt-5-mini's spread is the wider one.

The pilot said nine to twelve reps. I ran five.

Advised		12
	haiku 9, gpt-5-mini 12	
Run		5
	5 complete, a partial sixth	

5 run. 12 advised for a 5-point effect.

One posting, ten reps per model, before the main run.
Reps needed to resolve a five-point effect.

The twelve is an upper bound from one posting. The sixth rep was in progress when the run was stopped.

Each model returned more than one verdict across reps.

haiku-4-5		6
	6 of 21	
gpt-5-mini		12
	12 of 21	

6 of 21 postings, and 12 of 21.

Postings where a model's reps did not all return the same verdict, on the unedited resume.

Three verdict words, no rule in the prompt for when to use which. One resume.

A 22-point gap stands out. A two-point edit could hide.

EFFECT	FIVE REPS SAY
22.3, screener gap	well above noise
About 2, an edit	not ruled out
About 1, an edit	below the floor

Under about 3 points: not excluded.

My reading, anchored on the repo's sentence: effects under about three points are not excluded.

Twenty-one postings from eight employers treated as independent. The sentence goes next to every number.