

ZOEB NOMI

Bengaluru, India · +91 86008 16072 · zoeb.nomi@gmail.com

AI Product Manager · LLM Evaluation, Ground-Truth Data & RAG Quality

LinkedIn: <https://www.linkedin.com/in/zoebnomi> · Portfolio: <https://zoebnomi.com> · GitHub: <https://github.com/zoeb-nomi>

SUMMARY

AI Product Manager owning LLM output quality for a production tax-research and filing platform, where a wrong citation or a wrong line on a return is a compliance risk rather than a UX bug. Built the full quality loop behind the product: a self-maintaining 270K+-record primary-law corpus, human-expert golden datasets (Keystone benchmark; Project Evals - tax analysts authoring 2,000 reviewed datasets), and the evaluation layer that turns both into shipped fixes. Also built CrossSource, an open-source legal-RAG eval harness. Previously shipped a \$2.7M MRR enterprise product.

EXPERIENCE

Product Manager, Instead

Sept 2025 - Present

AI-native tax research, planning & filing platform; first new entrant in 20 years to clear IRS e-file approval alongside the incumbents - 2,540 federal, state, and local forms supported

- Own end-to-end output quality for production tax-research LLM, where wrong citations are a compliance risk: built the evaluation layer that makes quality measurable - flag to classify-by-failure-mode to weekly-review citation loop and a 4-dimension rubric (answerability, accuracy, citation quality, actionability) scored via LLM-as-judge with human validation - so regressions are caught before release, not by clients. Holds ~95% citation accuracy on the golden eval set.
- Grounded the model roadmap in evidence, not opinion: structured comparative experiments against three competing tax-AI research platforms (50+ scenarios, 5 rounds) and a "pipeline-collapse" eval testing whether the multi-step agent workflow (Tax Snapshot -> Strategies -> Memo -> QC) stays coherent and self-honest end-to-end - results set model-config and prompt decisions each round.
- Pressure-tested the platform against the public TaxCalcBench v2 suite - 1,200+ agent runs (50 cases, federal + 5 states) through a user-flow simulation harness across frontier models from four providers, permuting extended thinking, web search, and prompt style with full accuracy, latency, cost, and tool-call telemetry - to locate exactly where and why accuracy breaks; an LLM analyst-to-judge pipeline distilled the evidence into 7 shipped product improvements targeting those failure modes.
- Built Keystone, Instead's internal benchmark: a 75-case, 7-form golden set (1040/1120S/1065/1120/1041/990/5471 x 3 complexity tiers x federal/CA/NY), volume-weighted to real filing distribution, with its own eval harness and metric registry (953 scorecards;

strict/lenient scoring; per-tool telemetry across 27 tools). The baseline run located a "complexity cliff" that now sets the improvement roadmap.

- Designed and run Project Evals, Instead's human-expert ground-truth program: scenarios seeded from government e-file certification test cases (real-world-derived, PII-free) are worked end-to-end by in-house tax analysts simulating a firm-client engagement (~20 tasks per case: return, tax-saving plan, workbook, memo, client email), with tax-manager review as the acceptance gate - 2,000 human-authored, reviewed datasets targeted, benchmarking accuracy per task type and seeding model training.
- Built the retrieval foundation the product's answers stand on - and made it self-maintaining: a 270K+-record corpus of first-party primary law (federal and state statutes, regulations, case law and DoR rulings, plus local tax) spanning 230+ source types across all 51 jurisdictions, fed by a config-driven agentic ingestion pipeline (browser-agent collection to scraping to multi-model LLM cleaning and enrichment to MongoDB) that re-scrapes sources daily on update, so no one maintains the corpus by hand. Wrote the pipeline PRDs and schemas, built a Python fetchability harness across 378 sources, and led ~139K-document remediation - because retrieval quality caps answer quality.
- Eliminated the product's most dangerous failure mode - plausible-but-wrong citations that read as authoritative - by attaching document ID and sub-type to every cited source, and gave every strategy a legal backbone: a Legal Support Matrix across 160+ tax strategies (Tax Court cases, IRS rulings, audit guidance) surfaced through Source Explorer, command-palette search with cross-type related sources, so practitioners can verify any claim in seconds.
- Cut workflow latency ~90% (10-15 min to under a minute) by scoping incremental edits to changed inputs only - turning a batch process into an interactive tool practitioners iterate with - and hardened mid-workflow model-availability handling.
- Mapped Instead's e-file coverage against competitors across 6,201 forms in every US jurisdiction (2,540 supported), giving the filing roadmap a per-jurisdiction gap map.

Product Manager, Multiplier

Mar 2025 - Sept 2025

Global employment platform enabling compliant hiring across 150+ countries

- Owned the Value-Added Services vertical (procurement, BGV, ITSM, partnerships) and turned it into a revenue line: \$100K incremental revenue, a 47% efficiency gain from new SOPs, and the PosterElite compliance partnership launched end-to-end in 21 days.

Product Manager, Keka HR - 3 roles, 2 promotions

Dec 2022 - Mar 2025

India's leading HR technology platform; \$25M+ ARR CoreHR suite

- Built the Background Verification module zero-to-one (Checkr API) to open Keka's first international market: onboarded 28 US enterprise clients and \$2.7M MRR, anchoring the company's US expansion.

- Shipped 6 SSO integrations (Azure AD / Google Workspace) to remove enterprise adoption blockers - a 27% adoption lift and \$121K upsell; led the Exit Module revamp, cutting offboarding time 28% and lifting CSAT 80%.
- Ran GTM enablement at scale - product training and client webinars (4,400+ attendees) generating 500+ upsell requests and a 30% feature-adoption lift across CoreHR modules - because shipped features only count when clients adopt them.

Business Analyst, Hurix Digital

Oct 2021 - Oct 2022

- Led agile transformation of LMS development (43% faster delivery); drove 31% growth in customer conversions via a video-learning rollout.

OPEN-SOURCE PROJECTS

CrossSource - eval harness for citation-level reliability of legal RAG

<https://github.com/zoeb-nomi/crosssource>

- Built to prove eval craft in the open, on a domain where citations carry professional consequences: MIT-licensed and reproducible - 25 human-authored golden questions over 22 public court opinions; citation precision/recall with a failure taxonomy, faithfulness, and answer relevance via an LLM judge blind-validated at 100% human agreement (human labeling also caught a harness bug the judge missed). Citation-discipline prompting lifted precision 0.981 -> 0.994; recall (0.760) is retrieval-capped, not generation-capped.

SKILLS & TOOLS

AI / LLM: LLM evaluation, eval design, evaluation frameworks, agent evaluation, agentic workflows, LLM-as-a-judge, golden datasets, ground-truth & dataset design, human-expert annotation programs, human-in-the-loop, PII-safe data strategy, RAG evaluation, retrieval quality, citation precision/recall, failure-mode analysis, model benchmarking, data-pipeline automation, prompt engineering, latency/cost optimization, Python eval harnesses, MCP

Product: AI product management, product strategy, zero-to-one, enterprise B2B SaaS, roadmapping, API & platform PM, GTM, data-driven prioritization

Tools: Python, Linear, Figma, Metabase, Pendo, MongoDB, Claude / Claude Code, n8n, Cursor, Git

EDUCATION

STOA - General Management Program (2022)

MIT Aurangabad - B.E., Mechanical Engineering (2019)