

SUMMARY

AI Product Manager owning LLM output quality for a production tax-research and filing platform, where a wrong citation or a wrong line on a return is a compliance risk rather than a UX bug. Built the full quality loop behind the product: a self-maintaining 270K+-record primary-law corpus, human-expert golden datasets (Keystone benchmark; Project Evals - tax analysts authoring 2,000 reviewed datasets), & the evaluation layer that turns both into shipped fixes. Also built **CrossSource**, an open-source legal-RAG eval harness. Previously shipped a \$2.7M MRR enterprise product.

EXPERIENCE

Product Manager, Instead

Sept 2025 – Present

AI-native tax research, planning & filing platform; first new entrant in 20 years to clear IRS e-file approval alongside the incumbents - 2,540 federal, state, and local forms supported

- **Own end-to-end output quality for production tax-research LLM**, where wrong citations are a compliance risk: built the evaluation layer that makes quality measurable - flag → classify-by-failure-mode → weekly-review citation loop & 4-dimension rubric (answerability, accuracy, citation quality, actionability) scored via LLM-as-judge with human validation - so regressions are caught before release, not by clients. Holds ~95% citation accuracy on the golden eval set.
- **Grounded the model roadmap in evidence, not opinion**: structured comparative experiments against three competing tax-AI research platforms (50+ scenarios, 5 rounds) and a “pipeline-collapse” eval testing whether the multi-step agent workflow (Tax Snapshot → Strategies → Memo → QC) stays coherent and self-honest end-to-end - results set model-config and prompt decisions each round.
- **Pressure-tested the platform against the public TaxCalcBench v2 suite - 1,200+ agent runs** (50 cases, federal + 5 states) through a user-flow simulation harness across frontier models from four providers, permuting extended thinking, web search, and prompt style with full accuracy, latency, cost, and tool-call telemetry - to locate exactly where and why accuracy breaks; an LLM analyst→judge pipeline distilled the evidence into 7 shipped product improvements targeting those failure modes.
- **Built Keystone, Instead's internal benchmark**: a 75-case, 7-form golden set (1040/1120S/1065/1120/1041/990/5471 × 3 complexity tiers × federal/CA/NY), volume-weighted to real filing distribution, with its own eval harness and metric registry (953 scorecards; strict/lenient scoring; per-tool telemetry across 27 tools). The baseline run located a “complexity cliff” that now sets the improvement roadmap.
- **Designed and run Project Evals, Instead's human-expert ground-truth program**: scenarios seeded from government e-file certification test cases (real-world-derived, PII-free) are worked end-to-end by in-house tax analysts simulating a firm-client engagement (~20 tasks per case: return, tax-saving plan, workbook, memo, client email), with tax-manager review as the acceptance gate - 2,000 human-authored, reviewed datasets targeted, benchmarking accuracy per task type and seeding model training.
- **Built the retrieval foundation the product's answers stand on - and made it self-maintaining**: a 270K+-record corpus of first-party primary law (federal and state statutes, regulations, case law and DoR rulings, plus local tax) spanning 230+ source types across all 51 jurisdictions, fed by a config-driven agentic ingestion pipeline (browser-agent collection → scraping → multi-model LLM cleaning and enrichment → MongoDB) that re-scrapes sources daily on update, so no one maintains the corpus by hand. Wrote the pipeline PRDs and schemas, built a Python fetchability harness across 378 sources, and led ~139K-document remediation - because retrieval quality caps answer quality.
- **Eliminated the product's most dangerous failure mode** - plausible-but-wrong citations that read as authoritative - by attaching document ID and sub-type to every cited source, and gave every strategy a legal backbone: a Legal Support Matrix across 160+ tax strategies (Tax Court cases, IRS rulings, audit guidance) surfaced through Source Explorer, command-palette search with cross-type related sources, so practitioners can verify any claim in seconds.
- **Cut workflow latency ~90%** (10–15 min → under a minute) by scoping incremental edits to changed inputs only - turning a batch process into an interactive tool practitioners iterate with - & hardened mid-workflow model-availability handling.
- **Mapped Instead's e-file coverage against competitors across 6,201 forms in every US jurisdiction** (2,540 supported), giving the filing roadmap a per-jurisdiction gap map.

Product Manager, Multiplier

Mar 2025 – Sept 2025

Global employment platform enabling compliant hiring across 150+ countries

- **Owned the Value-Added Services vertical** (procurement, BGV, ITSM, partnerships) and turned it into a revenue line: \$100K incremental revenue, a 47% efficiency gain from new SOPs, and the PosterElite compliance partnership launched end-to-end in 21 days.

Product Manager, Keka HR - 3 roles, 2 promotions

Dec 2022 – Mar 2025

India's leading HR technology platform; \$25M+ ARR CoreHR suite

- **Built the Background Verification module zero-to-one** (Checkr API) to open Keka's first international market: onboarded 28 US enterprise clients and \$2.7M MRR, anchoring the company's US expansion.
- **Shipped 6 SSO integrations** (Azure AD / Google Workspace) to remove enterprise adoption blockers - a 27% adoption lift and \$121K upsell; led the Exit Module revamp, cutting offboarding time 28% and lifting CSAT 80%.
- **Ran GTM enablement at scale** - product training and client webinars (4,400+ attendees) generating 500+ upsell requests and a 30% feature-adoption lift across CoreHR modules - because shipped features only count when clients adopt them.

Business Analyst, Hurix Digital

Oct 2021 – Oct 2022

- **Led agile transformation of LMS development** (43% faster delivery); drove 31% growth in customer conversions via a video-learning rollout.

OPEN-SOURCE PROJECTS

CrossSource - eval harness for citation-level reliability of legal RAG

github.com/zoeb-nomi/crosssource

- Built to prove eval craft in the open, on a domain where citations carry professional consequences: MIT-licensed and reproducible - 25 human-authored golden questions over 22 public court opinions; citation precision/recall with a failure taxonomy, faithfulness, and answer relevance via an LLM judge blind-validated at 100% human agreement (human labeling also caught a harness bug the judge missed). Citation-discipline prompting lifted precision 0.981 → 0.994; recall (0.760) is retrieval-capped, not generation-capped.

SKILLS & TOOLS

AI / LLM: LLM evaluation · eval design · evaluation frameworks · agent evaluation · agentic workflows · LLM-as-a-judge · golden datasets · ground-truth & dataset design · human-expert annotation programs · human-in-the-loop · PII-safe data strategy · RAG evaluation · retrieval quality · citation precision/recall · failure-mode analysis · model benchmarking · data-pipeline automation · prompt engineering · latency/cost optimization · Python eval harnesses · MCP

Product: AI product management · product strategy · zero-to-one · enterprise B2B SaaS · roadmapping · API & platform PM · GTM · data-driven prioritization

Tools: Python · Linear · Figma · Metabase · Pendo · MongoDB · Claude / Claude Code · n8n · Cursor · Git

EDUCATION

STOA - General Management Program (2022) · MIT Aurangabad - B.E., Mechanical Engineering (2019)